

Homework #7

2024-10-09

Question 10.1 Using the same crime data set `uscrime.txt` as in Questions 8.2 and 9.1, find the best model you can using (a) a regression tree model, and (b) a random forest model. In R, you can use the `tree` package or the `rpart` package, and the `randomForest` package. For each model, describe one or two qualitative takeaways you get from analyzing the results (i.e., don't just stop when you have a good model, but interpret it too)

```
#housekeeping
library(pacman)
pacman::p_load(rio, tree, randomForest, caret, car, GGally, pROC)

crime <- import("D:/.../uscrime.txt")
str(crime)

## 'data.frame':  47 obs. of  16 variables:
## $ M      : num  15.1 14.3 14.2 13.6 14.1 12.1 12.7 13.1 15.7 14 ...
## $ So      : int   1 0 1 0 0 0 1 1 1 0 ...
## $ Ed      : num   9.1 11.3 8.9 12.1 12.1 11 11.1 10.9 9 11.8 ...
## $ Po1     : num   5.8 10.3 4.5 14.9 10.9 11.8 8.2 11.5 6.5 7.1 ...
## $ Po2     : num   5.6 9.5 4.4 14.1 10.1 11.5 7.9 10.9 6.2 6.8 ...
## $ LF      : num   0.51 0.583 0.533 0.577 0.591 0.547 0.519 0.542 0.553 0.632
## ...
## $ M.F     : num   95 101.2 96.9 99.4 98.5 ...
## $ Pop     : int   33 13 18 157 18 25 4 50 39 7 ...
## $ NW      : num   30.1 10.2 21.9 8 3 4.4 13.9 17.9 28.6 1.5 ...
## $ U1      : num   0.108 0.096 0.094 0.102 0.091 0.084 0.097 0.079 0.081 0.1
## ...
## $ U2      : num   4.1 3.6 3.3 3.9 2 2.9 3.8 3.5 2.8 2.4 ...
## $ Wealth: int  3940 5570 3180 6730 5780 6890 6200 4720 4210 5260 ...
## $ Ineq    : num   26.1 19.4 25 16.7 17.4 12.6 16.8 20.6 23.9 17.4 ...
## $ Prob    : num   0.0846 0.0296 0.0834 0.0158 0.0414 ...
## $ Time    : num   26.2 25.3 24.3 29.9 21.3 ...
## $ Crime   : int   791 1635 578 1969 1234 682 963 1555 856 705 ...

summary(crime)

##           M              So              Ed              Po1
## Min.      :11.90    Min.      :0.0000    Min.      : 8.70    Min.      : 4.50
## 1st Qu.:13.00    1st Qu.:0.0000    1st Qu.: 9.75    1st Qu.: 6.25
## Median :13.60    Median :0.0000    Median :10.80    Median : 7.80
## Mean      :13.86    Mean      :0.3404    Mean      :10.56    Mean      : 8.50
## 3rd Qu.:14.60    3rd Qu.:1.0000    3rd Qu.:11.45    3rd Qu.:10.45
## Max.      :17.70    Max.      :1.0000    Max.      :12.20    Max.      :16.60
##           Po2              LF              M.F              Pop
## Min.      : 4.100    Min.      :0.4800    Min.      : 93.40    Min.      : 3.00
```

## 1st Qu.: 5.850	1st Qu.:0.5305	1st Qu.: 96.45	1st Qu.: 10.00
## Median : 7.300	Median :0.5600	Median : 97.70	Median : 25.00
## Mean : 8.023	Mean :0.5612	Mean : 98.30	Mean : 36.62
## 3rd Qu.: 9.700	3rd Qu.:0.5930	3rd Qu.: 99.20	3rd Qu.: 41.50
## Max. :15.700	Max. :0.6410	Max. :107.10	Max. :168.00
## NW	U1	U2	Wealth
## Min. : 0.20	Min. :0.07000	Min. :2.000	Min. :2880
## 1st Qu.: 2.40	1st Qu.:0.08050	1st Qu.:2.750	1st Qu.:4595
## Median : 7.60	Median :0.09200	Median :3.400	Median :5370
## Mean :10.11	Mean :0.09547	Mean :3.398	Mean :5254
## 3rd Qu.:13.25	3rd Qu.:0.10400	3rd Qu.:3.850	3rd Qu.:5915
## Max. :42.30	Max. :0.14200	Max. :5.800	Max. :6890
## Ineq	Prob	Time	Crime
## Min. :12.60	Min. :0.00690	Min. :12.20	Min. : 342.0
## 1st Qu.:16.55	1st Qu.:0.03270	1st Qu.:21.60	1st Qu.: 658.5
## Median :17.60	Median :0.04210	Median :25.80	Median : 831.0
## Mean :19.40	Mean :0.04709	Mean :26.60	Mean : 905.1
## 3rd Qu.:22.75	3rd Qu.:0.05445	3rd Qu.:30.45	3rd Qu.:1057.5
## Max. :27.60	Max. :0.11980	Max. :44.00	Max. :1993.0

The data seems to have a wide range of scales—some variables are as small as the hundredths decimal place, while others, including the response variable, are in the thousands.

There is no missing data, but potential outliers might be present in variables like POP (upper end), NW (lower end), Prob (lower end), and Crime (top two values).

Next, using the tree library, carry out a tree regression, and use cross-validation to determine the ideal tree size.

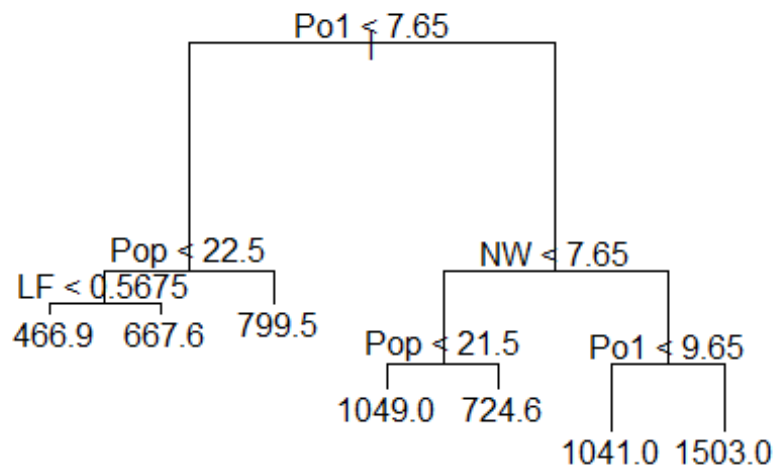
```
crimetree<-tree(Crime~.,data=crime)
summary(crimetree)

##
## Regression tree:
## tree(formula = Crime ~ ., data = crime)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "LF" "NW"
## Number of terminal nodes: 7
## Residual mean deviance: 47390 = 1896000 / 40
## Distribution of residuals:
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
## -573.900 -98.300  -1.545    0.000  110.600  490.100
```

There are 7 terminal nodes, with a residual mean deviance of 47,390.

Let's visualize the tree to examine the interactions and assess whether the splits are logical.

```
plot(crimetree)
text(crimetree ,pretty =0)
```



```

crimetree

## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 47 6881000  905.1
##    2) Po1 < 7.65 23  779200  669.6
##      4) Pop < 22.5 12  243800  550.5
##        8) LF < 0.5675 7  48520  466.9 *
##        9) LF > 0.5675 5  77760  667.6 *
##      5) Pop > 22.5 11  179500  799.5 *
##    3) Po1 > 7.65 24 3604000 1131.0
##      6) NW < 7.65 10  557600  886.9
##        12) Pop < 21.5 5  146400 1049.0 *
##        13) Pop > 21.5 5  147800  724.6 *
##      7) NW > 7.65 14 2027000 1305.0
##        14) Po1 < 9.65 6  170800 1041.0 *
##        15) Po1 > 9.65 8 1125000 1503.0 *

```

Po1 is identified as the most significant variable, which is why it is used for the initial split.

Now, let's compute both the R^2 and adjusted R^2 values for the tree model to evaluate its performance.

```

SST<-sum((crime$Crime-mean(crime$Crime))^2)
SSE<-sum((crime$Crime-predict(crimetree,newdata=crime))^2)
SST

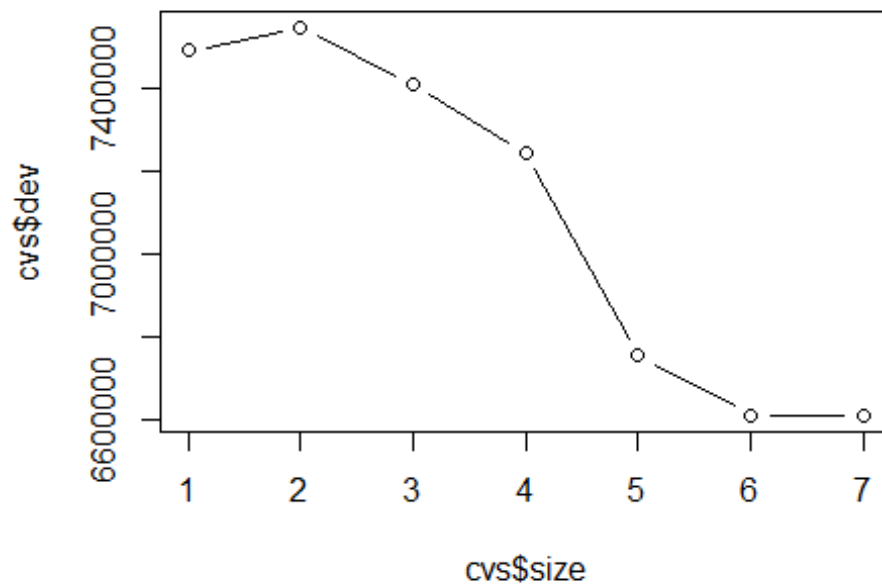
```

```
## [1] 6880928
SSE
## [1] 1895722
r2<-1-SSE/SST
r2
## [1] 0.7244962
n=nrow(crime)
k=7#number of terminal nodes
adjR2<-1-(1-r2)*(n-1)/(n-k-1)
adjR2
## [1] 0.6750468
```

Sum of Square Total: 6880928 Sum of Square Error: 1895722 R^2 : 0.7244962 Adjusted R^2 : 0.6750468

Let's try building a more refined tree model, aiming to reduce the sum of squared residuals and improve overall performance.

```
sample(1:1000,1)
## [1] 562
set.seed(123)
cvs<-cv.tree(crimetree, FUN = prune.tree, K=5)
cvs
## $size
## [1] 7 6 5 4 3 2 1
##
## $dev
## [1] 6609119 6609119 6752943 7245720 7410677 7547625 7491120
##
## $k
## [1] -Inf 117534.9 263412.9 355961.8 731412.1 1019362.7 2497521.7
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
plot(cvs$size ,cvs$dev ,type="b")
```



A 5-fold cross-validation shows that the optimal tree size is 2, as it results in the lowest error, with a deviance of 6,609,119.

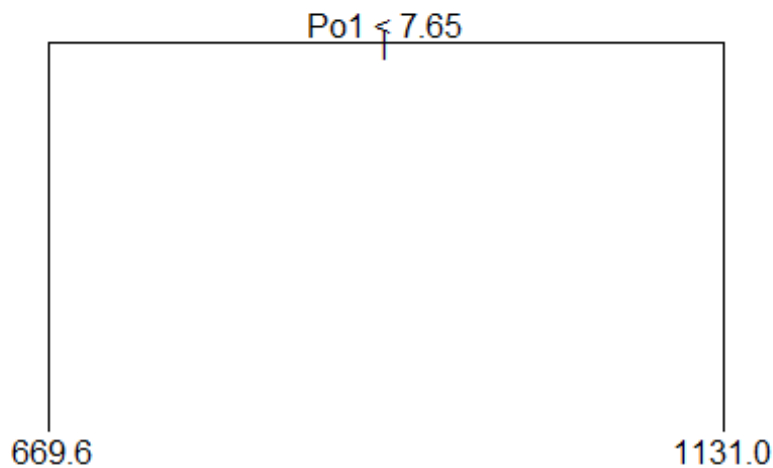
Since we now know that the optimal number of terminal nodes is 2, we can set this as a limit in the `tree.control` option to prevent overfitting if the tree exceeds 2 terminal nodes.

To create the recommended pruned tree, we should set the minimum number of observations at each node to 23. Let's rebuild the tree with `mincut = 23`, which will give us the desired 2 terminal nodes.

```
crimetree2 <- tree(Crime ~ ., data = crime, control = tree.control(nobs = 47, mincut = 23))
crimetree2

## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 47 6881000  905.1
##   2) Po1 < 7.65 23  779200  669.6 *
##   3) Po1 > 7.65 24 3604000 1131.0 *

plot(crimetree2)
text(crimetree2, pretty = 0)
```



This approach results in a tree with 2 terminal nodes (or leaves).

Now, let's evaluate the in-sample performance of the pruned tree to see how well it fits the training data.

```
summary(crimetree2)

##
## Regression tree:
## tree(formula = Crime ~ ., data = crime, control = tree.control(nobs = 47,
##   mincut = 23))
## Variables actually used in tree construction:
## [1] "Po1"
## Number of terminal nodes: 2
## Residual mean deviance: 97410 = 4383000 / 45
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -622.800 -193.200  -5.609    0.000  147.300   862.200
```

The in-sample residual mean deviance has now increased to 97,410, which is double what it was with 7 terminal nodes. This is expected, as the simpler model has fewer splits and thus explains less of the variance in the training data, leading to a higher residual deviance.

Let's now calculate the R^2 value for the pruned model to assess its explanatory power.

```
SSE2<-sum((crime$Crime-predict(crimetree2,newdata=crime))^2)
SSE2
```

```
## [1] 4383406

r2<-1-SSE2/SST
r2

## [1] 0.3629629

n=nrow(crime)
k=2#number of terminal nodes
adjR2<-1-(1-r2)*(n-1)/(n-k-1)
adjR2

## [1] 0.3340067
```

Sum of Square Error: 4383406 R^2 : 0.3629629 Adjusted R^2 : 0.3340067

The in-sample adjusted R^2 has dropped significantly, as expected for a simpler model. However, we anticipate that the cross-validated R^2 will be lower than that of the full model.

Next, we'll use the randomForest package to build a random forest model. For the number of variables to split on at each node (mtry), we will use the recommended value of $\sqrt{15}$, which rounds to 4. Let's proceed with mtry = 4 to train the random forest model.

```
library(randomForest)
rm<-randomForest(Crime~., data=crime, mtry=4)
rm

##
## Call:
## randomForest(formula = Crime ~ ., data = crime, mtry = 4)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 86008.71
##              % Var explained: 41.25
```

We observe that the random forest model's R^2 is 41.25%, which is an improvement over the selected tree model's R^2 of 36.2%.

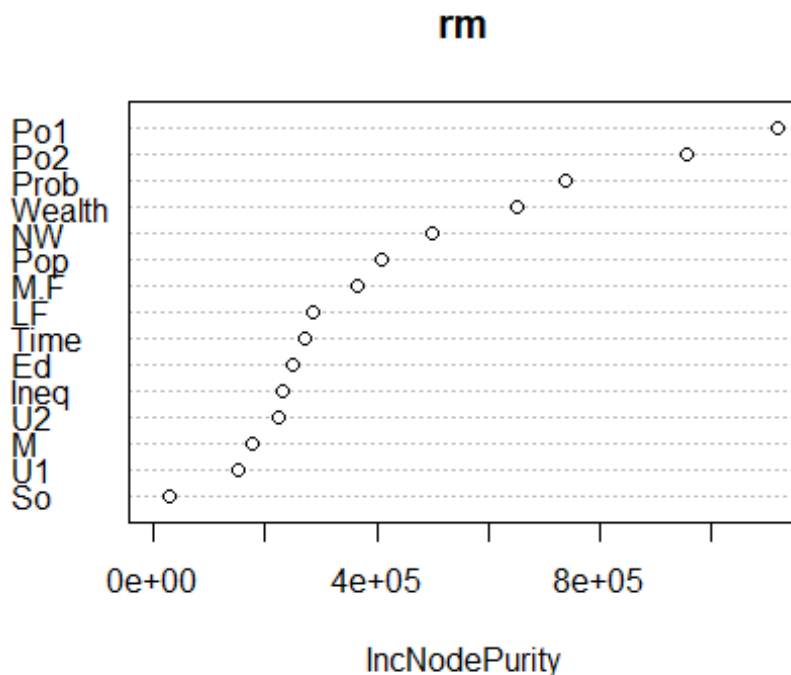
Next, let's evaluate the importance of the predictors using the Importance function and visualize it with VarImpPlot to identify the most influential variables in the model.

```
importance(rm)

##           IncNodePurity
## M           175616.94
## So           29094.31
## Ed           247982.57
## Po1          1117782.90
## Po2           956521.24
## LF           284317.91
```

```
## M.F      366642.57
## Pop      409715.40
## NW       500903.62
## U1       152274.19
## U2       222715.63
## Wealth   650623.69
## Ineq     230891.71
## Prob     739138.31
## Time     271831.04
```

```
varImpPlot(rm)
```



In Random Forest regression, the increase in node purity represents the average increase in RSS across all trees from splitting on a given variable. Once again, we find that Po1 is the most important variable.

Now, let's apply cross-validation to test different mtry levels and see if we can improve the model's performance further. By tuning mtry, we aim to find the optimal number of variables to split at each node, potentially boosting the model's accuracy.

```
cvrn<-train(Crime~., data=crime,
            trControl=trainControl(method = 'cv', number = 5, savePredictions
= TRUE),
            tuneGrid=expand.grid(mtry=1:15))
cvrm
```



```

## Random Forest
##
## 47 samples
## 15 predictors
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 38, 38, 37, 37, 38
## Resampling results across tuning parameters:
##
##   mtry  RMSE      Rsquared  MAE
##   1     297.6471  0.4652406  231.4440
##   2     282.9032  0.5010611  216.1343
##   3     283.1173  0.4807975  217.4543
##   4     283.9964  0.4744617  215.1611
##   5     285.3789  0.4640639  216.1700
##   6     288.9224  0.4495705  218.3318
##   7     294.8726  0.4486273  222.2781
##   8     291.2056  0.4440580  217.0495
##   9     296.8684  0.4332951  221.0979
##  10     300.0878  0.4218781  224.8335
##  11     303.5372  0.4145929  226.3279
##  12     296.8611  0.4288516  221.5051
##  13     302.1244  0.4221979  225.4605
##  14     301.0151  0.4124426  226.3083
##  15     304.2244  0.4037450  227.9303
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 2.

SSerm<-sum((cvrm$pred[,2]-cvrm$pred[,1])^2)
SSerm

## [1] 62272516

SSTrm<-sum((cvrm$pred[,2]-mean(cvrm$pred[,2]))^2)
SSTrm

## [1] 103213915

r2<-1-(SSerm/SSTrm)
r2

## [1] 0.3966655

```

Now, let's try reducing the number of trees built from 500 to 400 and compare the cross-validated performance. This could help prevent overfitting while maintaining model accuracy.

```

cvrm2<-train(Crime~., data=crime,
              trControl=trainControl(method = 'cv', number = 5, savePredictions

```

```

= TRUE),
  tuneGrid=expand.grid(mtry=1:15), ntree=400)
cvrm2

## Random Forest
##
## 47 samples
## 15 predictors
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 38, 38, 36, 37, 39
## Resampling results across tuning parameters:
##
##   mtry  RMSE      Rsquared  MAE
##   1    307.5985  0.5127369  234.3739
##   2    296.6861  0.5200015  226.0615
##   3    291.4671  0.5447273  223.0628
##   4    287.2255  0.5245885  217.5861
##   5    292.0182  0.5071527  219.0340
##   6    288.0024  0.5181279  214.9951
##   7    288.1204  0.5028579  213.7320
##   8    284.8795  0.5177826  212.3652
##   9    287.4861  0.5102370  215.0251
##  10    292.9740  0.5069761  218.3567
##  11    292.1732  0.5042289  218.0302
##  12    293.4688  0.4953211  218.1189
##  13    296.2153  0.4810546  220.7496
##  14    294.0484  0.5042943  217.9525
##  15    293.2615  0.4949142  216.6381
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 8.

SSerm<-sum((cvrm2$pred[,2]-cvrm2$pred[,1])^2)
SSerm

## [1] 64442310

SSTrm<-sum((cvrm2$pred[,2]-mean(cvrm2$pred[,2]))^2)
SSTrm

## [1] 103213915

r2<-1-(SSerm/SSTrm)
r2

## [1] 0.3756432

```

We've observed a slight improvement in cross-validated performance by reducing the number of trees to 400.

Let's now reduce the number of trees further, from 400 to 300, and compare the cross-validated performance to see if this helps prevent overfitting while maintaining or improving model accuracy.

```
set.seed(478)
cvrm3<-train(Crime~., data=crime,
             trControl=trainControl(method = 'cv', number = 5, savePredictions
= TRUE),
             tuneGrid=expand.grid(mtry=1:15), ntree=300)
cvrm3

## Random Forest
##
## 47 samples
## 15 predictors
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 36, 39, 35, 39, 39
## Resampling results across tuning parameters:
##
##   mtry  RMSE      Rsquared  MAE
##   1     291.9857  0.5106459  228.5653
##   2     275.7679  0.5424265  212.1414
##   3     270.5985  0.5616933  208.8549
##   4     268.0824  0.5812400  206.5469
##   5     267.5825  0.5652140  205.2438
##   6     266.1166  0.5824725  205.7723
##   7     269.1086  0.5597965  206.6419
##   8     264.8420  0.5825996  206.4396
##   9     267.0515  0.5680341  207.4476
##  10     268.2813  0.5673642  208.9184
##  11     274.4136  0.5351932  214.6021
##  12     268.8706  0.5578022  208.3199
##  13     270.7862  0.5528106  205.5508
##  14     274.7376  0.5360487  213.1964
##  15     269.8778  0.5659305  209.4186
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 8.

SSerm<-sum((cvrm3$pred[,2]-cvrm3$pred[,1])^2)
SSerm

## [1] 57281709

SSTrm<-sum((cvrm3$pred[,2]-mean(cvrm3$pred[,2]))^2)
SSTrm

## [1] 103213915
```

```
r2<-1-(SSErm/SSTrm)
r2
## [1] 0.4450195
```

Since reducing the number of trees to 300 didn't lead to any performance improvement, we'll stick with the earlier model, cvrm2, which was built using 400 trees. This model strikes a better balance between accuracy and overfitting.

Question 10.2 Describe a situation or problem from your job, everyday life, current events, etc., for which a logistic regression model would be appropriate. List some (up to 5) predictors that you might use.

A situation where a logistic regression model would be appropriate is predicting whether a customer will default on a loan (yes/no). Since the outcome is binary, logistic regression is a suitable approach.

5 potential predictors:

Income: The customer's annual income, which can help assess their ability to repay the loan.

Credit Score: A score indicating the customer's creditworthiness, an important factor in predicting loan defaults.

Debt-to-Income Ratio: The ratio of a customer's monthly debt payments to their monthly income.

Employment Status: Whether the customer is employed, unemployed, or self-employed, which can impact their financial stability.

Loan Amount: The total amount of the loan, as larger loans might be harder to repay depending on other factors.

Question 10.3.1 Using the GermanCredit data set `germancredit.txt` from <http://archive.ics.uci.edu/ml/machinelearning-databases/statlog/german/> (description at <http://archive.ics.uci.edu/ml/datasets/Statlog+%28German+Credit+Data%29>), use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not. Show your model (factors used and their coefficients), the software output, and the quality of fit. You can use the `glm` function in R. To get a logistic regression (logit) model on data where the response is either zero or one, use `family=binomial(link="logit")` in your `glm` function call.

```
credit<-import("D:/.../germancredit.txt")
str(credit)

## 'data.frame':    1000 obs. of  21 variables:
## $ V1 : chr  "A11" "A12" "A14" "A11" ...
## $ V2 : int   6 48 12 42 24 36 24 36 12 30 ...
## $ V3 : chr  "A34" "A32" "A34" "A32" ...
## $ V4 : chr  "A43" "A43" "A46" "A42" ...
```

```
## $ V5 : int 1169 5951 2096 7882 4870 9055 2835 6948 3059 5234 ...
## $ V6 : chr "A65" "A61" "A61" "A61" ...
## $ V7 : chr "A75" "A73" "A74" "A74" ...
## $ V8 : int 4 2 2 2 3 2 3 2 2 4 ...
## $ V9 : chr "A93" "A92" "A93" "A93" ...
## $ V10: chr "A101" "A101" "A101" "A103" ...
## $ V11: int 4 2 3 4 4 4 4 2 4 2 ...
## $ V12: chr "A121" "A121" "A121" "A122" ...
## $ V13: int 67 22 49 45 53 35 53 35 61 28 ...
## $ V14: chr "A143" "A143" "A143" "A143" ...
## $ V15: chr "A152" "A152" "A152" "A153" ...
## $ V16: int 2 1 1 1 2 1 1 1 1 2 ...
## $ V17: chr "A173" "A173" "A172" "A173" ...
## $ V18: int 1 1 2 2 2 2 1 1 1 1 ...
## $ V19: chr "A192" "A191" "A191" "A191" ...
## $ V20: chr "A201" "A201" "A201" "A201" ...
## $ V21: int 1 2 1 1 2 1 1 1 1 2 ...
```

summary(credit)

```
##          V1                V2                V3                V4
## Length:1000          Min.   : 4.0      Length:1000      Length:1000
## Class :character      1st Qu.:12.0      Class :character      Class :character
## Mode  :character      Median :18.0      Mode  :character      Mode  :character
##                               Mean   :20.9
##                               3rd Qu.:24.0
##                               Max.   :72.0
##          V5                V6                V7                V8
## Min.   : 250      Length:1000      Length:1000      Min.   :1.000
## 1st Qu.: 1366      Class :character      Class :character      1st Qu.:2.000
## Median : 2320      Mode  :character      Mode  :character      Median :3.000
## Mean   : 3271
## 3rd Qu.: 3972
## Max.   :18424
##          V9                V10               V11               V12
## Length:1000      Length:1000      Min.   :1.000      Length:1000
## Class :character      Class :character      1st Qu.:2.000      Class :character
## Mode  :character      Mode  :character      Median :3.000      Mode  :character
##                               Mean   :2.845
##                               3rd Qu.:4.000
##                               Max.   :4.000
##          V13               V14               V15               V16
## Min.   :19.00      Length:1000      Length:1000      Min.   :1.000
## 1st Qu.:27.00      Class :character      Class :character      1st Qu.:1.000
## Median :33.00      Mode  :character      Mode  :character      Median :1.000
## Mean   :35.55
## 3rd Qu.:42.00
## Max.   :75.00
##          V17               V18               V19               V20
## Length:1000          Min.   :1.000      Length:1000      Length:1000
```

```
## Class :character 1st Qu.:1.000 Class :character Class :character
## Mode :character Median :1.000 Mode :character Mode :character
## Mean :1.155
## 3rd Qu.:1.000
## Max. :2.000
## V21
## Min. :1.0
## 1st Qu.:1.0
## Median :1.0
## Mean :1.3
## 3rd Qu.:2.0
## Max. :2.0
```

One of the variables, V5 (credit amount), operates on a very different scale compared to the other variables. Thankfully, since we are using regression, this difference in scale won't disrupt the model. Additionally, there are no missing values in the data.

We will relabel the response variable to 0 and 1, where 0 represents 'Good' and 1 represents 'Bad,' with 'Bad' being treated as the positive class.

```
credit$V21<-ifelse(credit$V21==1,0,ifelse(credit$V21==2,1,credit$V21))
summary(credit$V21)

## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.0 0.0 0.0 0.3 1.0 1.0

prop.table(table(credit$V21))

##
## 0 1
## 0.7 0.3
```

The data shows that 70% of the cases are good credit risks, while 30% are bad.

Next, we'll split the dataset into 70% for training (roughly 700 observations) and 30% for testing. To ensure the splits maintain the same proportions of the response variable (good/bad credit risks), we will use stratified sampling. Additionally, the data will be randomized before splitting to avoid any ordering bias.

```
head(credit)

## V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13 V14 V15 V16 V17
## 1 A11 6 A34 A43 1169 A65 A75 4 A93 A101 4 A121 67 A143 A152 2 A173
## 2 A12 48 A32 A43 5951 A61 A73 2 A92 A101 2 A121 22 A143 A152 1 A173
## 3 A14 12 A34 A46 2096 A61 A74 2 A93 A101 3 A121 49 A143 A152 1 A172
## 4 A11 42 A32 A42 7882 A61 A74 2 A93 A103 4 A122 45 A143 A153 1 A173
```

```
## 5 A11 24 A33 A40 4870 A61 A73 3 A93 A101 4 A124 53 A143 A153 2 A173
2
## 6 A14 36 A32 A46 9055 A65 A73 2 A93 A101 4 A124 35 A143 A153 1 A172
2
##      V19 V20 V21
## 1 A192 A201 0
## 2 A191 A201 1
## 3 A191 A201 0
## 4 A191 A201 0
## 5 A191 A201 1
## 6 A192 A201 0

set.seed(1)
credit2<-credit[sample(1:nrow(credit)),]
head(credit2)

##      V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13 V14 V15 V16
V17 V18
## 836 A11 12 A30 A40 1082 A61 A73 4 A93 A101 4 A123 48 A141 A152 2
A173 1
## 679 A11 24 A32 A43 2384 A61 A75 4 A93 A101 4 A121 64 A141 A151 1
A172 1
## 129 A12 12 A34 A41 1860 A61 A71 4 A93 A101 2 A123 34 A143 A152 2
A174 1
## 930 A11 12 A33 A40 1344 A61 A73 4 A93 A101 2 A121 43 A143 A152 2
A172 2
## 509 A14 24 A32 A43 1413 A61 A73 4 A94 A101 2 A122 28 A143 A152 1
A173 1
## 471 A12 24 A32 A43 3092 A62 A72 3 A94 A101 2 A123 22 A143 A151 1
A173 1
##      V19 V20 V21
## 836 A191 A201 1
## 679 A191 A201 0
## 129 A192 A201 0
## 930 A191 A201 0
## 509 A191 A201 0
## 471 A192 A201 1
```

The dataset rows have now been randomized.

We can proceed by creating the building (training) set, ensuring that 70% of the data is allocated for model training while maintaining the stratified proportions of the response variable.

```
credit2$V21<-as.factor(credit2$V21)
sample(1:nrow(credit2),1)

## [1] 752

set.seed(266)
buildIndex <- createDataPartition(credit2$V21, p = .7, list=FALSE, times = 1)
```

```

build<-credit2[buildIndex,]
str(build)#confirm number of records is 70% of entire data set

## 'data.frame':    700 obs. of  21 variables:
## $ V1 : chr  "A11" "A11" "A12" "A14" ...
## $ V2 : int  12 24 12 24 18 18 30 24 15 24 ...
## $ V3 : chr  "A30" "A32" "A34" "A32" ...
## $ V4 : chr  "A40" "A43" "A41" "A43" ...
## $ V5 : int  1082 2384 1860 1413 2515 2427 4811 1442 874 3621 ...
## $ V6 : chr  "A61" "A61" "A61" "A61" ...
## $ V7 : chr  "A73" "A75" "A71" "A73" ...
## $ V8 : int   4 4 4 4 3 4 2 4 4 2 ...
## $ V9 : chr  "A93" "A93" "A93" "A94" ...
## $ V10: chr  "A101" "A101" "A101" "A101" ...
## $ V11: int   4 4 2 2 4 2 4 4 1 4 ...
## $ V12: chr  "A123" "A121" "A123" "A122" ...
## $ V13: int   48 64 34 28 43 42 24 23 24 31 ...
## $ V14: chr  "A141" "A141" "A143" "A143" ...
## $ V15: chr  "A152" "A151" "A152" "A152" ...
## $ V16: int   2 1 2 1 1 2 1 2 1 2 ...
## $ V17: chr  "A173" "A172" "A174" "A173" ...
## $ V18: int   1 1 1 1 1 1 1 1 1 1 ...
## $ V19: chr  "A191" "A191" "A192" "A191" ...
## $ V20: chr  "A201" "A201" "A201" "A201" ...
## $ V21: Factor w/ 2 levels "0","1": 2 1 1 1 1 1 1 2 1 2 ...

```

Now, let's verify if the proportion of 0s and 1s in the training set aligns with the original 70% "Good" (0) and 30% "Bad" (1) split observed in the full dataset. This check will confirm that the stratification worked correctly during the data splitting process.

```

prop.table(table(build$V21))

##
##    0    1
## 0.7 0.3

```

The stratified partition of the build set has been successfully created.

Now, we'll create a validation set using the remaining 30% of the data, ensuring that this set also maintains the same stratified proportions of the response variable (0s and 1s). This will be used for evaluating model performance.

```

test<-credit2[-buildIndex,]
str(test)#confirm number of records is 50% of test and validation data

## 'data.frame':    300 obs. of  21 variables:
## $ V1 : chr  "A11" "A12" "A14" "A12" ...
## $ V2 : int  12 24 24 9 18 6 6 12 27 18 ...
## $ V3 : chr  "A33" "A32" "A32" "A31" ...
## $ V4 : chr  "A40" "A43" "A43" "A41" ...
## $ V5 : int  1344 3092 999 5129 2404 368 1068 385 2570 1795 ...

```



```
## $ V6 : chr "A61" "A62" "A65" "A61" ...
## $ V7 : chr "A73" "A72" "A75" "A75" ...
## $ V8 : int 4 3 4 2 2 4 4 4 3 3 ...
## $ V9 : chr "A93" "A94" "A93" "A92" ...
## $ V10: chr "A101" "A101" "A101" "A101" ...
## $ V11: int 2 2 2 4 2 4 4 3 3 4 ...
## $ V12: chr "A121" "A123" "A123" "A124" ...
## $ V13: int 43 22 25 74 26 38 28 58 21 48 ...
## $ V14: chr "A143" "A143" "A143" "A141" ...
## $ V15: chr "A152" "A151" "A152" "A153" ...
## $ V16: int 2 1 2 1 2 1 1 4 1 2 ...
## $ V17: chr "A172" "A173" "A173" "A174" ...
## $ V18: int 2 1 1 2 1 1 2 1 1 1 ...
## $ V19: chr "A191" "A192" "A191" "A192" ...
## $ V20: chr "A201" "A201" "A201" "A201" ...
## $ V21: Factor w/ 2 levels "0","1": 1 2 1 2 1 1 1 1 2 1 ...
```

Let's now verify that the proportions of 0s and 1s in the test set match the original 70% "Good" (0) and 30% "Bad" (1) split observed in the full dataset. This will confirm that the stratification in the test set was applied correctly.

```
prop.table(table(test$V21))
```

```
##
##  0  1
## 0.7 0.3
```

The test set successfully matches the original 70-30% proportions.

Now, let's proceed with fitting a generalized linear model (GLM) using all the features on the build (training) set. This model will help us analyze the relationship between the predictors and the response variable.

```
glm1<-glm(V21~., data=build, family = binomial(link='logit'))
summary(glm1)

##
## Call:
## glm(formula = V21 ~ ., family = binomial(link = "logit"), data = build)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -6.093e-01  1.456e+00  -0.418  0.675662
## V1A12        -2.537e-01  2.695e-01  -0.941  0.346555
## V1A13        -7.945e-01  4.397e-01  -1.807  0.070779 .
## V1A14        -1.920e+00  2.984e-01  -6.434  1.24e-10 ***
## V2           4.573e-02  1.226e-02   3.731  0.000191 ***
## V3A31         3.289e-01  6.810e-01   0.483  0.629116
## V3A32        -6.354e-01  5.472e-01  -1.161  0.245570
## V3A33        -9.101e-01  6.051e-01  -1.504  0.132595
## V3A34        -1.522e+00  5.684e-01  -2.678  0.007404 **
```

```

## V4A41      -1.346e+00  4.672e-01  -2.881  0.003967 **
## V4A410     -1.405e+00  1.013e+00  -1.387  0.165519
## V4A42      -7.308e-01  3.308e-01  -2.209  0.027148 *
## V4A43      -1.162e+00  3.188e-01  -3.646  0.000266 ***
## V4A44      -9.327e-01  9.217e-01  -1.012  0.311606
## V4A45      -3.024e-01  7.040e-01  -0.430  0.667494
## V4A46       3.180e-01  5.185e-01   0.613  0.539609
## V4A48      -1.475e+01  5.199e+02  -0.028  0.977368
## V4A49      -8.913e-01  4.336e-01  -2.055  0.039838 *
## V5         9.189e-05  5.343e-05   1.720  0.085452 .
## V6A62      -4.411e-01  3.616e-01  -1.220  0.222523
## V6A63      -6.655e-01  5.573e-01  -1.194  0.232443
## V6A64      -2.053e+00  7.501e-01  -2.737  0.006191 **
## V6A65      -1.046e+00  3.363e-01  -3.110  0.001874 **
## V7A72       1.345e-01  5.498e-01   0.245  0.806785
## V7A73      -7.085e-02  5.304e-01  -0.134  0.893732
## V7A74      -7.997e-01  5.731e-01  -1.395  0.162876
## V7A75      -3.136e-01  5.248e-01  -0.597  0.550182
## V8         3.207e-01  1.076e-01   2.980  0.002886 **
## V9A92      -1.672e-01  4.735e-01  -0.353  0.723938
## V9A93      -6.027e-01  4.658e-01  -1.294  0.195685
## V9A94       1.628e-01  5.537e-01   0.294  0.768767
## V10A102     2.320e-01  4.790e-01   0.484  0.628179
## V10A103     -6.746e-01  5.182e-01  -1.302  0.192944
## V11        -5.896e-02  1.085e-01  -0.543  0.586956
## V12A122     1.673e-01  3.197e-01   0.524  0.600622
## V12A123     4.178e-01  2.898e-01   1.441  0.149444
## V12A124     7.226e-01  5.483e-01   1.318  0.187529
## V13        -2.583e-03  1.191e-02  -0.217  0.828265
## V14A142     2.878e-01  5.032e-01   0.572  0.567315
## V14A143     -4.091e-01  3.012e-01  -1.358  0.174345
## V15A152     -4.432e-01  2.858e-01  -1.551  0.120894
## V15A153     -9.597e-01  6.099e-01  -1.574  0.115597
## V16         3.717e-01  2.355e-01   1.578  0.114506
## V17A172     8.836e-01  8.888e-01   0.994  0.320133
## V17A173     8.121e-01  8.641e-01   0.940  0.347310
## V17A174     6.554e-01  8.719e-01   0.752  0.452252
## V18         3.294e-02  3.174e-01   0.104  0.917348
## V19A192     -2.985e-01  2.568e-01  -1.162  0.245071
## V20A202     -7.501e-01  6.976e-01  -1.075  0.282283
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 855.21  on 699  degrees of freedom
## Residual deviance: 590.96  on 651  degrees of freedom
## AIC: 688.96
##
## Number of Fisher Scoring iterations: 14

```

```
anova(glm1, test="Chi")#type I test

## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: V21
##
## Terms added sequentially (first to last)
##
##
```

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)	
## NULL			699	855.21		
## V1	3	126.472	696	728.74	< 2.2e-16	***
## V2	1	30.045	695	698.69	4.221e-08	***
## V3	4	19.919	691	678.77	0.0005182	***
## V4	9	24.864	682	653.91	0.0031251	**
## V5	1	0.389	681	653.52	0.5328840	
## V6	4	17.960	677	635.56	0.0012567	**
## V7	4	11.815	673	623.75	0.0187828	*
## V8	1	7.234	672	616.51	0.0071527	**
## V9	3	6.004	669	610.51	0.1114327	
## V10	2	2.838	667	607.67	0.2419908	
## V11	1	0.002	666	607.67	0.9626747	
## V12	3	2.567	663	605.10	0.4632270	
## V13	1	0.202	662	604.90	0.6527227	
## V14	2	3.977	660	600.92	0.1369079	
## V15	2	3.465	658	597.46	0.1768303	
## V16	1	2.322	657	595.13	0.1275229	
## V17	3	1.714	654	593.42	0.6338794	
## V18	1	0.007	653	593.41	0.9342698	
## V19	1	1.194	652	592.22	0.2745787	
## V20	1	1.258	651	590.96	0.2621147	

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Based on the Type I ANOVA test, we've identified several insignificant variables with p-values greater than 0.05, and we can drop the following: V20, V19, V18, V17, V15, V13, V12, V11, V10, V7, and V5.

Next, let's check for multicollinearity among the remaining predictors to ensure that no high correlations are distorting the model's estimates. This can be done by calculating the Variance Inflation Factor (VIF) for each variable to detect multicollinearity.

```
vif(glm1)
```

	GVIF	Df	GVIF^(1/(2*Df))
## V1	1.537492	3	1.074324
## V2	1.950812	1	1.396715
## V3	2.663751	4	1.130282
## V4	3.778548	9	1.076648

```
## V5 2.429089 1 1.558554
## V6 1.653463 4 1.064877
## V7 2.764675 4 1.135548
## V8 1.346400 1 1.160345
## V9 1.759161 3 1.098713
## V10 1.259315 2 1.059336
## V11 1.433695 1 1.197370
## V12 4.675313 3 1.293109
## V13 1.609002 1 1.268464
## V14 1.364883 2 1.080871
## V15 4.485618 2 1.455310
## V16 1.763185 1 1.327850
## V17 2.933300 3 1.196445
## V18 1.242080 1 1.114486
## V19 1.469455 1 1.212211
## V20 1.127144 1 1.061670
```

```
alias(glm1)
```

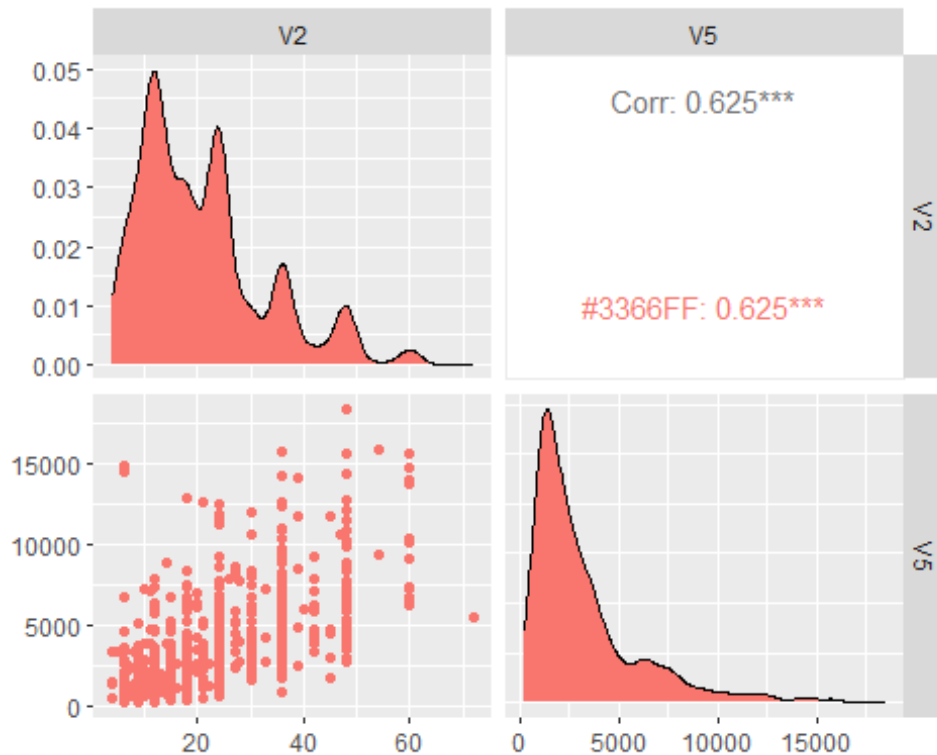
```
## Model :
```

```
## V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V7 + V8 + V9 + V10 + V11 +
##      V12 + V13 + V14 + V15 + V16 + V17 + V18 + V19 + V20
```

The Variance Inflation Factor (VIF) indicates that variables V2 (loan duration) and V5 (credit amount) have slightly elevated levels, suggesting potential multicollinearity. This could be due to a positive relationship between the duration of the loan and the loan amount, as larger loans may tend to have longer durations.

To explore this relationship further, we can create a scatter plot of V2 (loan duration) against V5 (credit amount) to visually assess the extent of their correlation. This will help us determine if there's a strong linear relationship between the two variables.

```
ggpairs(credit, columns = c('V2', 'V5'), mapping=ggplot2::aes(color=
'#3366FF'))
```



The scatter plot confirms a positive correlation between V2 (loan duration) and V5 (credit amount), and since the ANOVA test indicated that V5 can be dropped, we'll retain V2.

Now, we'll rerun the model, excluding the insignificant predictors: V20, V19, V18, V17, V15, V13, V12, V11, V10, V7, and V5. This will give us a more refined model based on the significant predictors.

```
glm2<-update(glm1, .~. -V20-V19-V18-V17-V15-V13-V12-V11-V10-V7-V5)
summary(glm2)
```

```
##
## Call:
## glm(formula = V21 ~ V1 + V2 + V3 + V4 + V6 + V8 + V9 + V14 +
##      V16, family = binomial(link = "logit"), data = build)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.321261   0.865875   0.371 0.710619
## V1A12         -0.224350   0.252630  -0.888 0.374509
## V1A13         -0.811500   0.421439  -1.926 0.054161 .
## V1A14        -1.888124   0.284197  -6.644 3.06e-11 ***
## V2             0.053450   0.009659   5.534 3.14e-08 ***
## V3A31         -0.037748   0.646336  -0.058 0.953427
## V3A32         -0.923920   0.521953  -1.770 0.076707 .
## V3A33         -1.157653   0.581363  -1.991 0.046451 *
## V3A34         -1.849084   0.544090  -3.398 0.000678 ***
## V4A41         -1.165924   0.434173  -2.685 0.007245 **
```

```
## V4A410      -1.391403    0.891563   -1.561 0.118610
## V4A42       -0.548035    0.307107   -1.785 0.074341 .
## V4A43       -1.186527    0.296972   -3.995 6.46e-05 ***
## V4A44       -0.782282    0.893089   -0.876 0.381069
## V4A45       -0.229949    0.647040   -0.355 0.722301
## V4A46        0.503749    0.496017    1.016 0.309825
## V4A48      -14.781920  545.259460   -0.027 0.978372
## V4A49       -0.950098    0.408022   -2.329 0.019883 *
## V6A62       -0.276723    0.334854   -0.826 0.408578
## V6A63       -0.592588    0.543414   -1.090 0.275497
## V6A64       -1.852736    0.714181   -2.594 0.009481 **
## V6A65       -1.001513    0.318204   -3.147 0.001647 **
## V8           0.221999    0.094778    2.342 0.019165 *
## V9A92       -0.212221    0.440190   -0.482 0.629726
## V9A93       -0.782688    0.431170   -1.815 0.069483 .
## V9A94       -0.025271    0.520422   -0.049 0.961270
## V14A142      0.418859    0.487310    0.860 0.390047
## V14A143     -0.386378    0.290892   -1.328 0.184095
## V16          0.300935    0.219125    1.373 0.169644
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 855.21  on 699  degrees of freedom
## Residual deviance: 616.50  on 671  degrees of freedom
## AIC: 674.5
##
## Number of Fisher Scoring iterations: 14
```

With the new model based on the selected terms, we've observed an improvement in the AIC, which has decreased to 674.5—indicating a better fit.

However, the 4th variable has an unusual coefficient estimate for level A48, and both levels A44 and A48 have very high p-values. We need to further investigate these levels to determine whether it makes sense to combine them with another level to simplify the model and improve its performance.

Let's explore their distribution and relationship with the response variable to make an informed decision on potential combinations.

[View\(build\\$V4\)](#)

After combining factor level A48 (loan purpose: retraining) with the base level A40 (car new), as well as combining A44 with the base due to low observation counts, let's rerun the model and check if the AIC improves. This adjustment should help stabilize the coefficient estimates and reduce any instability in the model.

```

levels(build$V4)[levels(build$V4)==c("A48", "A44")]<-"A40"
glm2.5<-update(glm2,~.)
summary(glm2.5)

##
## Call:
## glm(formula = V21 ~ V1 + V2 + V3 + V4 + V6 + V8 + V9 + V14 +
##      V16, family = binomial(link = "logit"), data = build)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.321261   0.865875   0.371 0.710619
## V1A12         -0.224350   0.252630  -0.888 0.374509
## V1A13         -0.811500   0.421439  -1.926 0.054161 .
## V1A14        -1.888124   0.284197  -6.644 3.06e-11 ***
## V2             0.053450   0.009659   5.534 3.14e-08 ***
## V3A31         -0.037748   0.646336  -0.058 0.953427
## V3A32         -0.923920   0.521953  -1.770 0.076707 .
## V3A33        -1.157653   0.581363  -1.991 0.046451 *
## V3A34        -1.849084   0.544090  -3.398 0.000678 ***
## V4A41         -1.165924   0.434173  -2.685 0.007245 **
## V4A410        -1.391403   0.891563  -1.561 0.118610
## V4A42         -0.548035   0.307107  -1.785 0.074341 .
## V4A43        -1.186527   0.296972  -3.995 6.46e-05 ***
## V4A44         -0.782282   0.893089  -0.876 0.381069
## V4A45         -0.229949   0.647040  -0.355 0.722301
## V4A46          0.503749   0.496017   1.016 0.309825
## V4A48        -14.781920  545.259460  -0.027 0.978372
## V4A49         -0.950098   0.408022  -2.329 0.019883 *
## V6A62         -0.276723   0.334854  -0.826 0.408578
## V6A63         -0.592588   0.543414  -1.090 0.275497
## V6A64        -1.852736   0.714181  -2.594 0.009481 **
## V6A65        -1.001513   0.318204  -3.147 0.001647 **
## V8             0.221999   0.094778   2.342 0.019165 *
## V9A92         -0.212221   0.440190  -0.482 0.629726
## V9A93         -0.782688   0.431170  -1.815 0.069483 .
## V9A94         -0.025271   0.520422  -0.049 0.961270
## V14A142        0.418859   0.487310   0.860 0.390047
## V14A143       -0.386378   0.290892  -1.328 0.184095
## V16           0.300935   0.219125   1.373 0.169644
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 855.21  on 699  degrees of freedom
## Residual deviance: 616.50  on 671  degrees of freedom
## AIC: 674.5
##
## Number of Fisher Scoring iterations: 14

```

Now, let's proceed with evaluating the cross-validated performance of the model using 7 folds, with each fold containing 100 observations. This will give us a more robust measure of the model's performance beyond the in-sample results.

```
build$risk<-as.factor(ifelse(build$V21==1,'bad','good'))
buildfolds<-createFolds(build$risk,k=7)

set.seed(123)
cvrm<-train(risk ~ V1 + V2 + V3 + V4 + V6 + V8 + V9 + V14 + V16,
            data=build,
            method='glm',
            trControl=trainControl(method = 'cv', number = 7, index =
buildfolds, classProbs = TRUE, summaryFunction = twoClassSummary),
            metric='ROC')

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```



```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful
cases
```

```
cvrn
```

```
## Generalized Linear Model
##
## 700 samples
## 9 predictor
## 2 classes: 'bad', 'good'
##
## No pre-processing
## Resampling: Cross-Validated (7 fold)
## Summary of sample sizes: 100, 100, 100, 100, 100, 100, ...
## Resampling results:
##
## ROC Sens Spec
## 0.6880943 0.5063492 0.7792517
```

```
cvrn$finalModel
```

```
##
## Call: NULL
##
## Coefficients:
## (Intercept) V1A12 V1A13 V1A14 V2
V3A31
## -0.32126 0.22435 0.81150 1.88812 -0.05345
0.03775
## V3A32 V3A33 V3A34 V4A41 V4A410
V4A42
## 0.92392 1.15765 1.84908 1.16592 1.39140
0.54804
## V4A43 V4A44 V4A45 V4A46 V4A48
V4A49
## 1.18653 0.78228 0.22995 -0.50375 14.78192
0.95010
```

```
##          V6A62          V6A63          V6A64          V6A65          V8
V9A92
##      0.27672      0.59259      1.85274      1.00151      -0.22200
0.21222
##          V9A93          V9A94          V14A142          V14A143          V16
##      0.78269      0.02527      -0.41886      0.38638      -0.30094
##
## Degrees of Freedom: 699 Total (i.e. Null); 671 Residual
## Null Deviance:      855.2
## Residual Deviance: 616.5      AIC: 674.5
```

The 7-fold cross-validated results show an area under the curve (AUC) of 0.6880943, sensitivity of 0.5063492, and specificity of 0.7792517 with a 50% threshold, assuming '1' as the positive class. As noted, some predicted probabilities were exactly 0 or 1, which can occur with logistic regression.

Now, let's make predictions on the test set after relabeling V4A48 and V4A44 to V4A40, just as we did in the build set, and evaluate the performance using the 50% threshold. This will help us see how well the model generalizes to unseen data.

```
levels(test$V4)[levels(test$V4)==c("A48","A44")]<-"A40"
glm2test<-predict(glm2.5, newdata=test,type='response')

#set threshold at 50%
glm2testfact<-as.factor(ifelse(glm2test>0.5,1,0))#positive class is bad risks

confusionMatrix(glm2testfact,as.factor(test$V21))#using 50% threshold

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 182  51
##              1  28  39
##
##              Accuracy : 0.7367
##              95% CI : (0.683, 0.7856)
##              No Information Rate : 0.7
##              P-Value [Acc > NIR] : 0.09179
##
##              Kappa : 0.3236
##
## Mcnemar's Test P-Value : 0.01332
##
##              Sensitivity : 0.8667
##              Specificity : 0.4333
##              Pos Pred Value : 0.7811
##              Neg Pred Value : 0.5821
##              Prevalence : 0.7000
```

```
##          Detection Rate : 0.6067
##    Detection Prevalence : 0.7767
##          Balanced Accuracy : 0.6500
##
##          'Positive' Class : 0
##

library(pROC)
roc(test$V21,ifelse(glm2test>0.5,1,0))#inputs must be numeric

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
##
## Call:
## roc.default(response = test$V21, predictor = ifelse(glm2test >      0.5, 1,
## 0))
##
## Data: ifelse(glm2test > 0.5, 1, 0) in 210 controls (test$V21 0) < 90 cases
## (test$V21 1).
## Area under the curve: 0.65
```

From the test set results, we can see that the model correctly predicted 58.21% of the bad risks, but it misclassified 51 out of 90 bad risks as good, which could lead to significant costs for the bank. On the other hand, 28 good risks were classified as bad, but this mistake has a lower cost since these risks were still good.

Test set sensitivity for the positive class ('1') is higher than the cross-validated result from the build set, which is beneficial since the cost of misclassifying bad risks as good is higher. However, test set specificity for the negative class ('0') is lower, but this is less concerning because the primary focus is on avoiding false negatives.

The test set area under the curve (AUC) is lower than the cross-validated AUC of 0.69, which suggests the model's performance on the test data isn't as strong as during cross-validation, but the sensitivity improvement offsets some of this concern, given the higher cost of misclassifying bad risks as good ones.

Question 10.3.2 Because the model gives a result between 0 and 1, it requires setting a threshold probability to separate between "good" and "bad" answers. In this data set, they estimate that incorrectly identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad. Determine a good threshold probability based on your model.

Let's explore other thresholds and apply our cost function to assess the model's performance.

We'll define the costs accordingly.

```
mycost=function(r,pi,threshold){#bad risks
```

```

    ifelse(r==1,ifelse(pi<threshold,5,0), #misclassifying bad risk incurs a
cost of 5 per risk

        ifelse(pi>threshold#misclass good risk incurrs a cost of 1 per risk
,1,0))
    }

    for(threshold in (1:100)/100){
cost= sum(mycost(test$V21,glm2test,threshold))
print(cbind(threshold,cost))}

##      threshold cost
## [1,]      0.01  209
##      threshold cost
## [1,]      0.02  202
##      threshold cost
## [1,]      0.03  202
##      threshold cost
## [1,]      0.04  194
##      threshold cost
## [1,]      0.05  193
##      threshold cost
## [1,]      0.06  195
##      threshold cost
## [1,]      0.07  194
##      threshold cost
## [1,]      0.08  202
##      threshold cost
## [1,]      0.09  204
##      threshold cost
## [1,]      0.1   212
##      threshold cost
## [1,]      0.11  212
##      threshold cost
## [1,]      0.12  207
##      threshold cost
## [1,]      0.13  204
##      threshold cost
## [1,]      0.14  202
##      threshold cost
## [1,]      0.15  197
##      threshold cost
## [1,]      0.16  197
##      threshold cost
## [1,]      0.17  198
##      threshold cost
## [1,]      0.18  197
##      threshold cost
## [1,]      0.19  195

```

```
##      threshold cost
## [1,]      0.2  201
##      threshold cost
## [1,]      0.21 205
##      threshold cost
## [1,]      0.22 200
##      threshold cost
## [1,]      0.23 206
##      threshold cost
## [1,]      0.24 211
##      threshold cost
## [1,]      0.25 211
##      threshold cost
## [1,]      0.26 207
##      threshold cost
## [1,]      0.27 205
##      threshold cost
## [1,]      0.28 200
##      threshold cost
## [1,]      0.29 215
##      threshold cost
## [1,]      0.3  214
##      threshold cost
## [1,]      0.31 218
##      threshold cost
## [1,]      0.32 217
##      threshold cost
## [1,]      0.33 216
##      threshold cost
## [1,]      0.34 231
##      threshold cost
## [1,]      0.35 229
##      threshold cost
## [1,]      0.36 232
##      threshold cost
## [1,]      0.37 236
##      threshold cost
## [1,]      0.38 243
##      threshold cost
## [1,]      0.39 250
##      threshold cost
## [1,]      0.4  253
##      threshold cost
## [1,]      0.41 252
##      threshold cost
## [1,]      0.42 253
##      threshold cost
## [1,]      0.43 256
##      threshold cost
## [1,]      0.44 265
```

```
##      threshold cost
## [1,]      0.45 269
##      threshold cost
## [1,]      0.46 277
##      threshold cost
## [1,]      0.47 275
##      threshold cost
## [1,]      0.48 279
##      threshold cost
## [1,]      0.49 278
##      threshold cost
## [1,]      0.5  283
##      threshold cost
## [1,]      0.51 288
##      threshold cost
## [1,]      0.52 297
##      threshold cost
## [1,]      0.53 300
##      threshold cost
## [1,]      0.54 303
##      threshold cost
## [1,]      0.55 308
##      threshold cost
## [1,]      0.56 318
##      threshold cost
## [1,]      0.57 322
##      threshold cost
## [1,]      0.58 321
##      threshold cost
## [1,]      0.59 325
##      threshold cost
## [1,]      0.6  330
##      threshold cost
## [1,]      0.61 338
##      threshold cost
## [1,]      0.62 341
##      threshold cost
## [1,]      0.63 339
##      threshold cost
## [1,]      0.64 339
##      threshold cost
## [1,]      0.65 339
##      threshold cost
## [1,]      0.66 353
##      threshold cost
## [1,]      0.67 365
##      threshold cost
## [1,]      0.68 369
##      threshold cost
## [1,]      0.69 374
```

```
##      threshold cost
## [1,]      0.7  379
##      threshold cost
## [1,]      0.71 378
##      threshold cost
## [1,]      0.72 377
##      threshold cost
## [1,]      0.73 376
##      threshold cost
## [1,]      0.74 381
##      threshold cost
## [1,]      0.75 380
##      threshold cost
## [1,]      0.76 385
##      threshold cost
## [1,]      0.77 385
##      threshold cost
## [1,]      0.78 384
##      threshold cost
## [1,]      0.79 384
##      threshold cost
## [1,]      0.8  384
##      threshold cost
## [1,]      0.81 394
##      threshold cost
## [1,]      0.82 404
##      threshold cost
## [1,]      0.83 414
##      threshold cost
## [1,]      0.84 414
##      threshold cost
## [1,]      0.85 418
##      threshold cost
## [1,]      0.86 423
##      threshold cost
## [1,]      0.87 427
##      threshold cost
## [1,]      0.88 432
##      threshold cost
## [1,]      0.89 436
##      threshold cost
## [1,]      0.9  436
##      threshold cost
## [1,]      0.91 435
##      threshold cost
## [1,]      0.92 445
##      threshold cost
## [1,]      0.93 445
##      threshold cost
## [1,]      0.94 445
```

```
##      threshold cost
## [1,]      0.95 450
##      threshold cost
## [1,]      0.96 450
##      threshold cost
## [1,]      0.97 450
##      threshold cost
## [1,]      0.98 450
##      threshold cost
## [1,]      0.99 450
##      threshold cost
## [1,]      1 450
```

The minimum cost is 193, which occurs at a threshold of 0.05.

Now, let's examine the confusion matrix at this threshold to evaluate the model's classification performance.

```
#set threshold at 5%
glm2testfact2<-as.factor(ifelse(glm2test>0.05,1,0))#positive class is bad risks

confusionMatrix(glm2testfact2,as.factor(test$V21))#using 5% threshold

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0  37   4
##              1 173  86
##
##              Accuracy : 0.41
##              95% CI : (0.3538, 0.468)
##      No Information Rate : 0.7
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.0857
##
##      McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.1762
##              Specificity : 0.9556
##              Pos Pred Value : 0.9024
##              Neg Pred Value : 0.3320
##              Prevalence : 0.7000
##              Detection Rate : 0.1233
##      Detection Prevalence : 0.1367
##              Balanced Accuracy : 0.5659
##
```



```
##          'Positive' Class : 0
##
roc(test$V21,ifelse(glm2test>0.05,1,0))#inputs must be numeric
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
##
## Call:
## roc.default(response = test$V21, predictor = ifelse(glm2test >      0.05,
1, 0))
##
## Data: ifelse(glm2test > 0.05, 1, 0) in 210 controls (test$V21 0) < 90
cases (test$V21 1).
## Area under the curve: 0.5659
```

Lowering the threshold to 0.05 reduced the number of false positives to 4. However, this also increased the number of false negatives to 173.

It's important to note that overall accuracy isn't the key metric here, due to the class imbalance (70% good, 30% bad) and the 5:1 cost ratio for misclassification. In fact, overall accuracy decreased, but this is expected given our focus on minimizing costs, not maximizing accuracy.